



§18a — 1BQF vs classical at small n (shots=819200, device=CPU, reps=3)


Segment-level fidelity



Track-level efficiency



Solution scatter at $n_{\text{trk}} = 6$, rep=0



Post-selection & walltime

