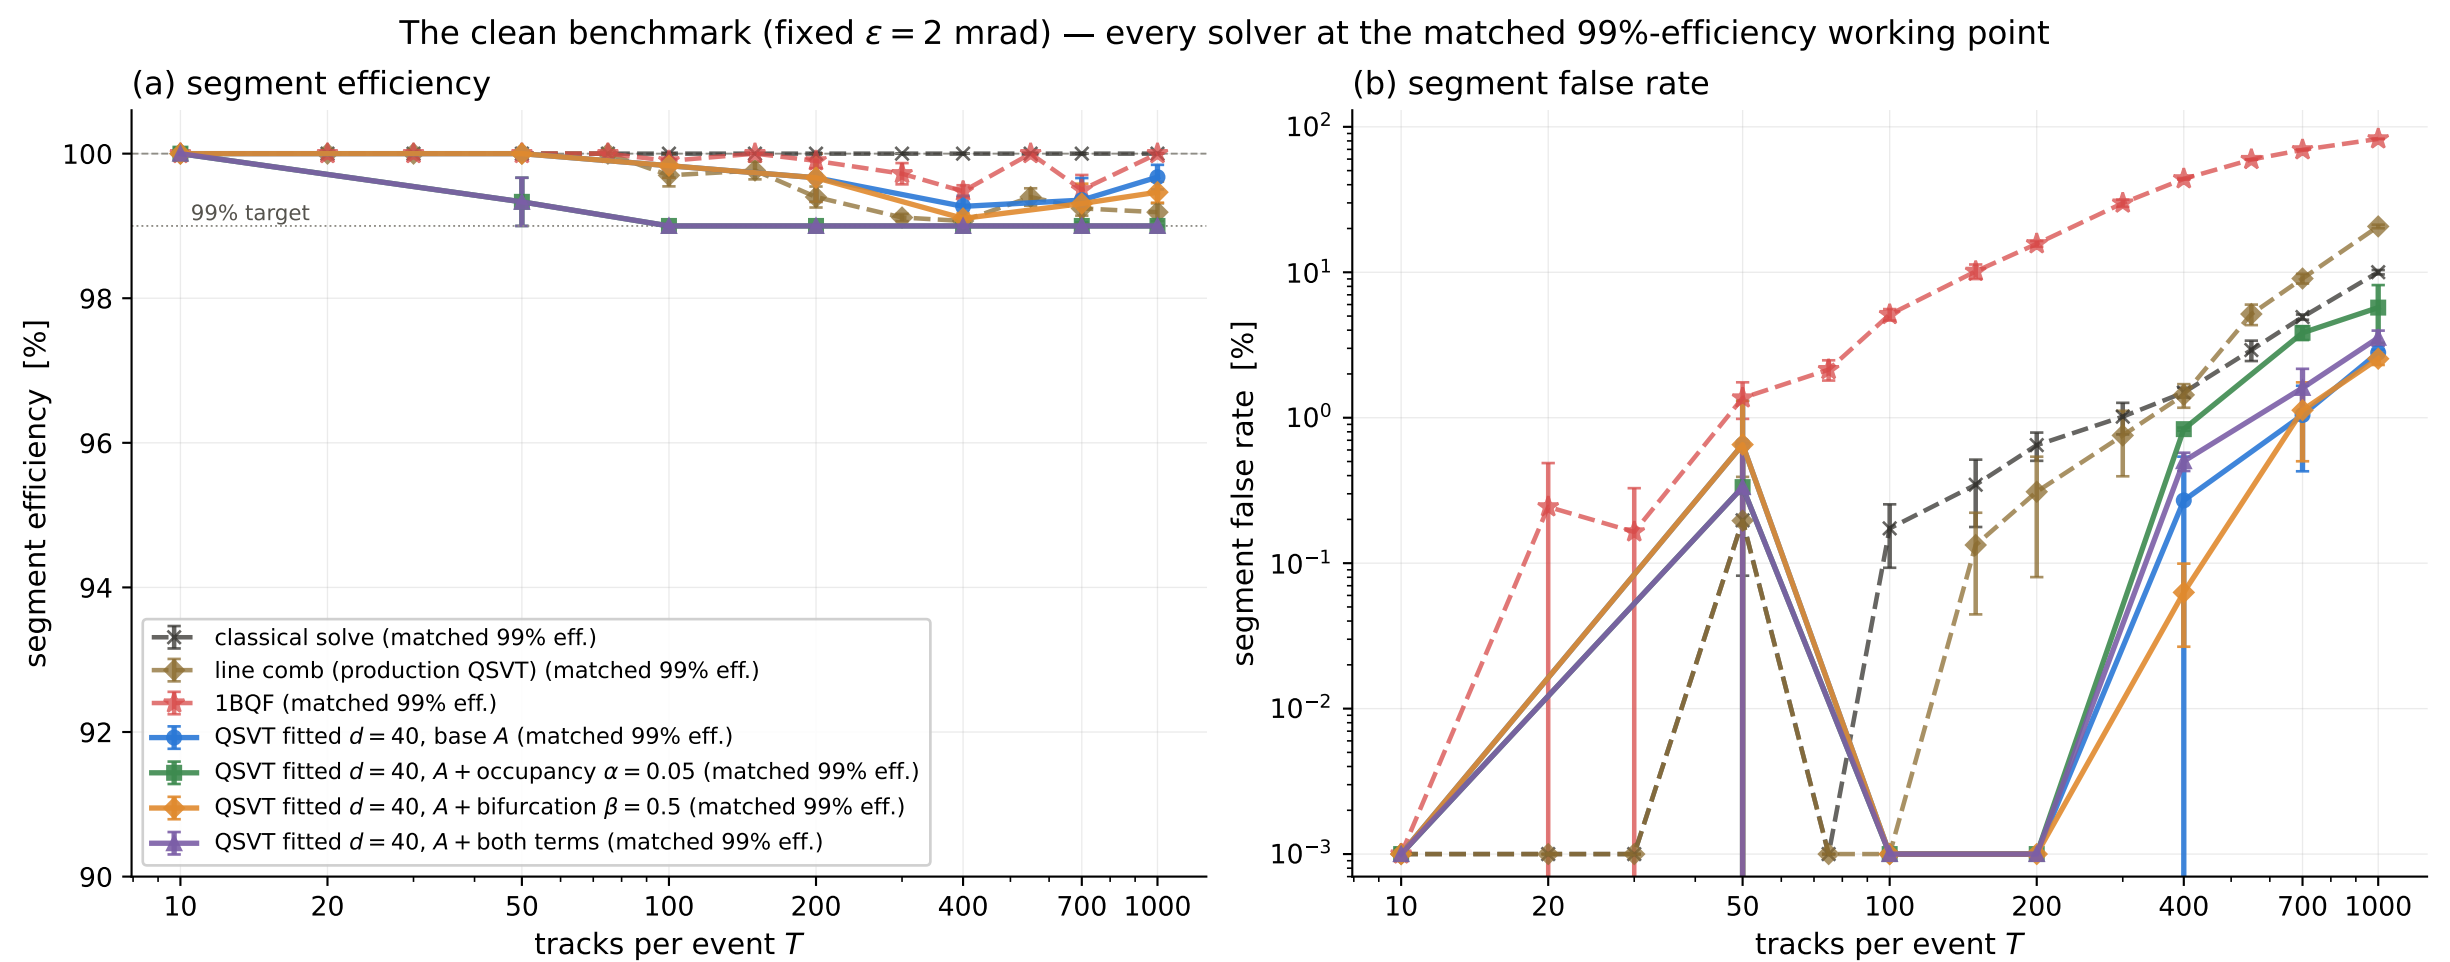

one configuration throughout: $\sigma_{\text{scatt}} = 10^{-4}$, $\sigma_{\text{res}} = 0$, no hit drop, $\varphi_{\max} = 0.2$, FIXED $\varepsilon = 2$ mrad, $\gamma = 3$, $\delta = 1$, step kernel · store curves: campaign mean±sem over reps (classical 5–20, line comb 5–10, 1BQF 3–10 per T) · fitted series: the same campaign events, reps 0–2, mean±sem, response refit per operator (moment-space ridge, $d = 40$) · the threshold of every curve is set per event so that 99% of the true segments survive; panel (a) shows the convention holding, panel (b) carries the comparison · a measured 0% false rate is drawn at the 10^{-3} floor